

## Networking continued

`ip addr show --info` on all interfaces

`ip -s link show eno1` --- get packet stats on an int. Int now not like eth0 but (more like bsd?) dependent on HW

`/etc/sysconfig/network-scripts`

Recommended: check out scripts in `/etc/sysconfig/network-scripts/ifcfg-*`

They are fairly readable and there is one for each connection/ interface

Hostname is in `/etc/hostname`

`hostnamectl set-hostname vm0.example.com` -- to set the hostname

`hostnamectl status` -- to request the hostname

DNS resolver is pushed from `/etc/sysconfig/network-scripts/ifcfg-*` to `/etc/resolv.conf`

## IPv6

`ip a` -----lists... same as `-a?` `ip link show`

`nmcli con add con-name blah type ethernet ifname eth0 ip6 2001:db8:0:10::d000:310/64 gw6 2001:db8:0:10::1 ipv4 192.168.1.5 gw4 192.168.1.1`

Remember, new connections use IP4 or IP6 with NO "v" - modifying (using a "+") DOES use a "v":

`nmcli con mod blah +ipv6.dns 2001:4860:4860::8888`

## Setting up routing

Rather than jump into `ip route` commands, this goes to `nmtui`

NMTUI opens with edit and activate (a connection) options and "set system hostname"

Routing - "edit a connection"

Under all of the standard IPv4 options is "Routing" with an edit hyperlink

It will say "no custom routes" and give you an "add" hyperlink

It asks for destination/prefix, next hop, and metric. click ok and exit

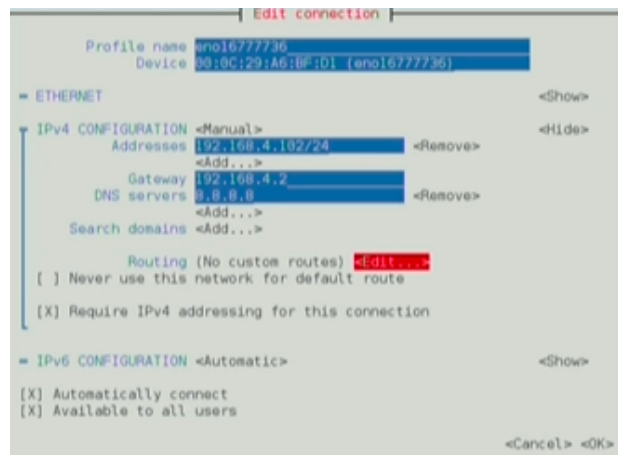
`ip route show` won't show it unless you reopen `nmtui` and click "activate connection"

It will show up in `/etc/sysconfig/network-scripts/` When you open it, all it has is this:

ADDRESS0:10.0.0.0

NETMASK0:255.255.255.0

GATEWAY0: 192.168.4.4



## Virtual Bridges

eno1 to virbr0 bridges vnet0 and vnet1 virtual network interfaces, and each of those has one separate host attached. Remember that these "2 networks" are merely switched - the explanation made it sound like 2 subnets

These will be messing with vish constructs:

`virsh list --all` --get a list of the VM runnings

`brctl show` -- we see virbr0, whether stp is enabled, the bridgeid, the interfaces that are up

ip link show --baremetal lo and eno1 listed, and then virbr, vnet0 and vnet1  
 on the vnet0 and 1 these values are added: pfifo\_fast master virbr0 state UNKNOWN mode default qlen 500  
 on virbr0, noqueue state UP mode DEFAULT  
 on eno1 - pfifo\_fast state UP mode DEFAULT  
 All of these have the standard other stuff: <multicast, broadcast, up, lower\_up>, mtu 1500 qdisk link/ether (mac address) brd ff:ff:ff:ff:ff:ff

### **The RHCE requirement is knowing how to set up a bridge**

yum install -y bridge-utils --- You need to install this, and you have to know the current interfaces:

nmcli dev show -- Get the id of the interface, eno0166777723 for our ethernet

nmcli dev disconnect eno0166777723

nmcli con add type bridge-slave con-name br0-port1 ifname eno0166777723 master br0

-- do this to add the interface we disconnected to the bridge- It will complain the bridge isn't there, we make it below

-- you repeat this for the other interfaces that will be part of this bridge, br0

-- if you only have one interface, thats fine, two you can do load balancing etc for HP

nmcli con add type bridge con-name br0 ifname br0 -- defines the bridge

brctl show --will show your bridges

-- interesting to note, the bridge id is 8000.000000000000

Here is what we get when we check out the files in /etc/sysconfig/network-scripts/

Scared you won't remember these? **Man pages! See nmcli-examples(5) page (is at the bottom of "man nmcli")**

```
[root@server2 network-scripts]# cat ifcfg-br0
DEVICE=br0
STP=yes
BRIDGING_OPTS=priority=32768
TYPE=Bridge
BOOTPROTO=dhcp
DEFROUTE=yes
PEERDNS=yes
PEERROUTES=yes
IPV4_FAILURE_FATAL=no
IPV6INIT=yes
IPV6_AUTOCONF=yes
IPV6_DEFROUTE=yes
IPV6_PEERDNS=yes
IPV6_PEERROUTES=yes
IPV6_FAILURE_FATAL=no
NAME=br0
UUID=da404f2b-f967-4d86-8bb5-bca4bfc44ee8
ONBOOT=yes
[root@server2 network-scripts]# cat ifcfg-br0-port1
TYPE=Ethernet
NAME=br0-port1
UUID=91ead89a-0bfb-46da-856b-e7b18ad723d2
DEVICE=eno16677736
ONBOOT=yes
BRIDGE=br0
[root@server2 network-scripts]#
```

The nmcli-examples man page has everything you need to type, as long as you know what you are doing!!

### **TeamD**

**Network Link Aggregation - Bonding is deprecated. Teaming takes it's place.**

Modes "runners": broadcast, roundrobin, activebackup (one line active, the other a backup), loadbalance, and LACP.

### **There are Four Steps to Network Teaming**

- Create team interface
  - nmcli con add type team con-name team0 ifname team0 config '{"runner": {"name": "loadbalance"}}'
- Determine network configuration
  - nmcli con mod team0 ipv4.addresses 10.0.0.10/24
  - nmcli con mod team0 ipv4.method manual
- Assign the port interfaces
  - nmcli con add type team-slave ifname eth0 master team0 con-name team0-eth0
  - nmcli con add type team-slave ifname eth1 master team0 con-name team0-eth1
- Bring team and port interfaces up/down
  - nmcli con up team0
  - nmcli dev dis eth0; nmcli dev dis eth1
- Verify: teamdctl team0 state

Note: Above, the order in "Assign port interfaces" is wrong (see example) con-name team0-eth0 goes before ifname!

### **Creating a Bridge Based on Network Teams**

- Doesn't work with NetworkManager enabled
- Modify the team configuration file ifcfg-team0 and add BRIDGE=brteam0 to tell it to connect to the bridge device
- Make sure no IP configuration remains in the ifcfg-team0-port files
- Manually create a bridge file:  
    DEVICE=brteam0  
    TYPE=Bridge  
    IPADDR0=192.168.122.100  
    PREFIX0=24
- Check **man 5 nmcli-examples!**

## Command line example- complete teamd and link aggregation setup using teaming

```
[root@server2 ~]# nmcli con add type team con-name team0 ifname team0 config '{"runner": {"name": "loadbalance"}}'
Connection 'team0' (1c62729f-8685-4174-ae56-8e4b7daec063) successfully added.
[root@server2 ~]# nmcli con mod team0 ipv4.addresses 10.0.0.10/24
[root@server2 ~]# nmcli con mod team0 ipv4.method manual
[root@server2 ~]# ip link show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eno16777736: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast master br0 state UP mode DEFAULT qlen 1000
    link/ether 00:0c:29:a6:bf:d1 brd ff:ff:ff:ff:ff:ff
3: eno33554992: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT qlen 1000
    link/ether 00:0c:29:a6:bf:d0 brd ff:ff:ff:ff:ff:ff
4: eno50332216: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc pfifo_fast state DOWN mode DEFAULT qlen 1000
    link/ether 00:0c:29:a6:bf:e5 brd ff:ff:ff:ff:ff:ff
5: br0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT
    link/ether 00:0c:29:a6:bf:d1 brd ff:ff:ff:ff:ff:ff
6: team0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT
    link/ether 3a:f6:cb:f4:ce:5c brd ff:ff:ff:ff:ff:ff
[root@server2 ~]# nmcli con add type team-slave ifname eno33554992 master team0 con-name team0-eno33554992
Error: Unexpected argument 'con-name'
[root@server2 ~]# man nmcli-examples
[root@server2 ~]# nmcli con add type team-slave con-name eno33554992 ifname eno33554992 master team0
Connection 'eno33554992' (4bcecd02-4031-44e7-8812-51efeba4e69b) successfully added.
[root@server2 ~]# nmcli con add type team-slave con-name eno50332216 ifname eno50332216 master team0
Connection 'eno50332216' (c7b310c5-411b-477b-856e-d14025878a16) successfully added.
[root@server2 ~]# nmcli con up team0
Connection successfully activated (D-Bus active path: /org/freedesktop/NetworkManager/ActiveConnection/4)
[root@server2 ~]# nmcli dev dis eno50332216
Error: Device 'eno50332216' [/org/freedesktop/NetworkManager/Devices/4] disconnecting failed: This device is not active
[root@server2 ~]# nmcli dev dis eno33554992
[root@server2 ~]# teamdctl team0 state
setup:
  runner: loadbalance
[root@server2 ~]# ip link show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eno16777736: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast master br0 state UP mode DEFAULT qlen 1000
    link/ether 00:0c:29:a6:bf:d1 brd ff:ff:ff:ff:ff:ff
3: eno33554992: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT qlen 1000
    link/ether 00:0c:29:a6:bf:d0 brd ff:ff:ff:ff:ff:ff
4: eno50332216: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc pfifo_fast state DOWN mode DEFAULT qlen 1000
    link/ether 00:0c:29:a6:bf:e5 brd ff:ff:ff:ff:ff:ff
5: br0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT
    link/ether 00:0c:29:a6:bf:d1 brd ff:ff:ff:ff:ff:ff
6: team0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT
    link/ether 3a:f6:cb:f4:ce:5c brd ff:ff:ff:ff:ff:ff
[root@server2 ~]#
```

### Exercise 6

- Shut down your virtual machine. If possible, at this point make a snapshot of the virtual machine to easily restore the original configuration after the exercise
  - Add a new virtual network card to the virtual machine
  - Add a network team interface and call it team0 and assign both network interfaces to the team
  - Put the team driver in a network bridge
- 
- ip link
  - nmcli dev dis eth0; nmcli dev dis eth1
  - nmcli con add type team con-name team0 ifname team0 config '{"runner": {"name": "activebackup"}}'
  - nmcli con add type team-slave con-name team0-port1 ifname eth0 master team0; repeat for eth1
  - teamdctl team0 state
  - At this point the team works. To put it in a bridge we need to disable NetworkManager and disable the team0 driver as the bridge will take the configuration instead of the team driver!
  - nmcli dev dis team0
  - systemctl stop NetworkManager; systemctl disable NetworkManager
  - yum install bridge-utils(!)
  - vim /etc/sysconfig/network-scripts/ifcfg-team0; add BRIDGE=brteam0
  - Remove the IP configuration from the ifcfg-team0-port files
  - Create the ifcfg-brteam0 file with following contents:  
DEVICE=brteam0  
ONBOOT=yes  
TYPE=Bridge  
IPADDR0=192.168.122.100  
PREFIX0=24
  - systemctl restart network